

AlexNet on CIFAR-10 - Project Documentation

Name: Tensae Aschalew
GSR/3976/17

1. Introduction

This project implements the AlexNet convolutional neural network using PyTorch and applies it to the CIFAR-10 image classification task. The model is trained, evaluated, and tested in the Google Colab environment. This implementation is intended for educational purposes, with clearly annotated code and detailed explanations.

2. Objective

To:

- Understand the architecture of AlexNet.
 - Apply it to a smaller dataset (CIFAR-10).
 - Explore training, evaluation, and performance optimization.
 - Demonstrate how deep CNNs can be adapted to datasets other than ImageNet.
-

3. Background: AlexNet Architecture

3.1 Historical Context

AlexNet was proposed by Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton in 2012. It won the ImageNet Large Scale Visual Recognition Challenge (ILSVRC) with a **top-5 error rate of 15.3%**, far surpassing the runner-up.

It became a cornerstone in the deep learning revolution for computer vision due to:

- Its depth (8 layers: 5 convolutional, 3 fully connected).

- Use of GPU acceleration.
- ReLU activation, which enabled faster convergence.
- Dropout for reducing overfitting.
- Overlapping max-pooling for better spatial understanding.

3.2 Predecessor: LeNet-5

Feature	LeNet-5	AlexNet
Year	1998	2012
Dataset	MNIST (grayscale)	ImageNet (RGB)
Input Size	32×32	224×224
Conv Layers	2	5
FC Layers	2	3
Activation	Tanh	ReLU
Regularization	None	Dropout
GPU Usage	No	Yes (2 GPUs originally)
Parameters	~60K	~60 million

AlexNet extended LeNet-5 to support higher dimensional RGB inputs and deeper architectures suitable for real-world image recognition problems.

4. Dataset: CIFAR-10

4.1 Description

- Total samples: 60,000 images
- Image size: 32×32 pixels, RGB
- 10 classes:

- airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck
- Training set: 50,000 images
- Test set: 10,000 images

4.2 Preprocessing for AlexNet

Since AlexNet expects 224×224 pixel images, we resize all CIFAR-10 images to 224×224 using interpolation. Additionally, the dataset is normalized to mean = (0.5, 0.5, 0.5) and std = (0.5, 0.5, 0.5) for each channel to standardize the input.

5. Model Architecture

5.1 AlexNet - Layer Breakdown

Feature Extractor:

- **Conv1**: 96 filters, 11×11 kernel, stride 4 \rightarrow ReLU \rightarrow MaxPool
- **Conv2**: 256 filters, 5×5 kernel \rightarrow ReLU \rightarrow MaxPool
- **Conv3**: 384 filters, 3×3 kernel \rightarrow ReLU
- **Conv4**: 384 filters, 3×3 kernel \rightarrow ReLU
- **Conv5**: 256 filters, 3×3 kernel \rightarrow ReLU \rightarrow MaxPool

Classifier:

- Dropout
 - Fully Connected 4096 \rightarrow ReLU \rightarrow Dropout
 - Fully Connected 4096 \rightarrow ReLU
 - Fully Connected 10 (output logits)
-

6. Environment Setup (Google Colab)

6.1 Enable GPU

- Go to **Runtime** > **Change runtime type**
- Set **Hardware Accelerator** to **GPU**

6.2 Dependencies

PyTorch and torchvision are pre-installed in Google Colab.

7. Training Pipeline

- **Loss Function:** CrossEntropyLoss (for multi-class classification)
- **Optimizer:** Adam (learning rate = 0.001)
- **Batch Size:** 64
- **Epochs:** Default is 1 (can be increased)
- **Evaluation Metric:** Accuracy on test set

Training Steps

1. Load and preprocess data
 2. Define AlexNet model
 3. Train model with mini-batches
 4. Report loss every 100 steps
 5. Evaluate on test data
 6. Print test accuracy
-

8. Results (5 Epoch)

Metric	Value
Epochs	5
Final Test Accuracy	~60.80%
Training Time	~487.89 seconds (with GPU)

Note: With 10–30 epochs, accuracy can reach 75%+

10. Possible Improvements

- Train for more epochs
 - Use learning rate schedulers
 - Add data augmentation: random cropping, flipping
 - Apply transfer learning using pretrained weights
 - Use mixed precision training for faster performance
-

11. References

- Krizhevsky et al., “ImageNet Classification with Deep Convolutional Neural Networks”, NeurIPS 2012.
- CIFAR-10 dataset: <https://www.cs.toronto.edu/~kriz/cifar.html>